

University of Warsaw

Faculty of Psychology

Kamil Tomaszek

Record book number: 432044

Middle Polish Dependency Treebank in Universal Dependencies Format: Design, Implementation, and Analysis

Master's thesis
in COGNITIVE SCIENCE

The thesis was written under the supervision of

Dr. Alina Wróblewska

Institute of Computer Science

Polish Academy of Sciences

Dr. Grzegorz Krajewski

University of Warsaw

Warsaw, November 2025

Summary

This thesis presents a rule-based approach to converting the Middle Polish Dependency Treebank (MPDT), annotated in a Polish-specific scheme, into the Universal Dependencies (UD) format. After introducing the project motivation, data sources, and target standard, the thesis outlines general design assumptions behind the conversion, the mapping strategy, and the validation workflow. It reports overall outcomes of the conversion and sketches applications and extensions, including releasing MPDT-UD and implications for research in historical language processing within cognitive science.

Keywords

Middle Polish, dependency trees, treebank conversion, Universal Dependencies

The title of the thesis in Polish

Średniopolski Bank Drzew Zależnościowych w formacie Universal Dependencies: projekt, implementacja i analiza

Streszczenie

Praca przedstawia podejście regułowe do konwersji Średniopolskiego Banku Drzew Zależnościowych (MPDT), anotowanego w polskim schemacie, do formatu Universal Dependencies (UD). Po krótkim omówieniu motywacji, danych i standardu docelowego zaprezentowano ogólne założenia projektu, strategię odwzorowań oraz schemat walidacji. Przedstawiono ogólne wyniki konwersji oraz możliwe zastosowania i kierunki rozwoju, w tym udostępnienie MPDT-UD i znaczenie dla badań nad przetwarzaniem języka historycznego w kognitywistyce.

Słowa kluczowe

język średniopolski, drzewa zależnościowe, konwersja korpusu, Universal Dependencies

The title of the thesis in English

Middle Polish Dependency Treebank in Universal Dependencies Format: Design, Implementation, and Analysis

Contents

1. Introduction	6
1.1. Motivation	6
1.2. Objectives	7
1.3. Contributions	7
1.4. Structure of the Thesis	7
2. Background	9
2.1. Dependency Grammar	9
2.2. Universal Dependencies	11
2.3. Middle Polish Linguistic Resources	15
2.3.1. KorBa	15
2.3.2. MPDT	16
3. Linguistic Features of Middle Polish	19
3.1. Orthography and Punctuation	19
3.1.1. Orthography and Transliteration	19
3.1.2. Punctuation	20
3.2. Morphology	21
3.2.1. Additional Parts of Speech and Morphological Features	21
3.2.2. Gender System and Declension	23
3.3. Syntax	23
3.3.1. Word Order and Non-projectivity	23
3.3.2. Predicate Ellipsis	24
3.3.3. Clause Linking and Subordination	25
3.4. Summary	26
4. Conversion Design and Implementation	27
4.1. Converter Architecture and Environment	27
4.1.1. Core Data Structures	27
4.1.2. Project Repository Structure	28
4.2. Conversion Pipeline	29

4.3.	Phase 1: Morphosyntactic Conversion	30
4.3.1.	Pre-conversion	30
4.3.2.	Core POS Conversion	30
4.3.3.	Post-conversion	31
4.4.	Phase 2: Dependency Conversion	31
4.4.1.	Structural Restructuring	31
4.4.2.	Label Mapping	35
4.4.3.	Correction and Post-processing	35
4.5.	Audibility and Processing Workflow	36
4.5.1.	Logging and Traceability	36
4.5.2.	Processing Workflow	36
4.6.	Summary	36
5.	Validation and Outcomes	38
5.1.	Validation Workflow and Formal Conformance	38
5.1.1.	Iterative Validation Procedure	38
5.1.2.	Validation Dataset	39
5.1.3.	Official UD Validator	39
5.2.	Conformance Results	40
5.2.1.	Handling Idiosyncratic Edge Cases	40
5.3.	Outcomes: the MPDT-UD 1.0 Treebank	41
5.3.1.	Corpus Size and Data Split	41
5.3.2.	Release, Licensing, and UD Integration	42
6.	Applications and Cognitive Science Perspective	43
6.1.	Usefulness and Audience	43
6.1.1.	Who Benefits and How	44
6.1.2.	Packaging and License of the Converter	44
6.1.3.	Repository and Reproducible UD Integration	45
6.2.	Use Cases	45
6.2.1.	Historical Syntax and Diachrony	46
6.2.2.	Parser Training and Evaluation	46
6.3.	Cognitive Science Perspective	47
6.3.1.	Processing Constraints	47
6.3.2.	Category Change Over Time	48
6.4.	Future Work	49
6.4.1.	Coverage and Phenomena	49
6.4.2.	Generalization and Automation	50

Chapter 1

Introduction

1.1. Motivation

Natural-language preprocessing tools and comparative treebank research have standardized around Universal Dependencies (UD), which enables typologically informed analyses and cross-lingual transfer (Nivre et al. 2020). For Polish texts from the 17th and 18th centuries, however, key resources remain outside UD: texts in the KorBa corpus (Gruszczyński et al. 2022) and the emerging Middle Polish Dependency Treebank (MPDT) are annotated in a Polish-specific scheme (Wieczorek 2025). KorBa is a corpus of historical Polish texts, while MPDT adds a syntactic dependency layer to selected portions of this corpus. However, these resources being annotated in a different format creates challenges for interoperability with UD-based tools and limits straightforward comparative studies with other languages.

A natural solution is to convert these resources to the UD format. From an engineering perspective, however, a faithful, auditable conversion is non-trivial: historical orthography, abbreviations, clitic mobility, numeral complexes, and multiword conjunctions/prepositions interact with head rules and label inventories. Prior conversion experience for contemporary Polish offers valuable guidance (Wróblewska 2018; Wróblewska 2020), yet historical data introduce additional phenomena that require explicit, rule-based handling and transparent traceability.

As Wieczorek (2025) notes, MPDT’s current format is well-suited to comparative studies with contemporary Polish syntax; at the same time, she highlights the advantages of moving to UD for cross-linguistic comparability, wider intelligibility, and representational options such as enhanced dependencies for shared dependents and shared governors in coordination—even if some information may be lost in conversion.

This thesis operationalizes that rationale by delivering a documented, UD-oriented converter for MPDT and preparing the current version of MPDT-UD suitable for validation and downstream use. The intended users include historical linguists needing UD-compatible data and NLP practitioners interested in diachronic Polish or cross-

lingual experiments.

1.2. Objectives

The thesis pursues the following research goals:

- (R1) **Design a UD-oriented conversion strategy for MPDT.** Specify mapping principles that respect Middle Polish specifics while aligning with UD guidelines.
- (R2) **Implement an auditable conversion pipeline.** Provide modular components for morphosyntax mapping and dependency restructuring, with token-level logging.
- (R3) **Ensure UD conformance and evaluability.** Produce output that passes the official UD validator (on all levels) and supports downstream analysis.
- (R4) **Document decisions.** Record non-obvious mapping choices and edge-case policies to enable maintenance and reuse.

1.3. Contributions

This project delivers concrete, reusable artifacts:

- (C1) **A rule-based MPDT → MPDT-UD converter.** A modular pipeline with fine-grained logging, selectively adapting ideas from Wróblewska (2018) while targeting Middle Polish phenomena. The code will be released in a public repository under an open-source license, together with this thesis, which documents the design and implementation.
- (C2) **An initial public release of MPDT-UD.** A set of MPDT (2018 sentences at the time of writing) converted automatically and validated with the official UD validator.

1.4. Structure of the Thesis

- **Chapter 2: Background.** Introduces dependency grammar and the Polish Dependency Bank (PDB) annotation scheme; outlines the Universal Dependencies (UD) framework, including its layers and relation inventories; and presents the key source resources—KorBa and MPDT—that the conversion operates on.

- **Chapter 3: Linguistic Features of Middle Polish.** Describes linguistic properties of Middle Polish relevant to conversion: orthography and punctuation practices; characteristic morphological categories (`adjb`, `ppasb`, `ppraet`, dual number); evolving masculine gender distinctions; the connective *jako* in its comparative and role uses; and syntactic features such as non-projectivity and predicate ellipsis.
- **Chapter 4: Conversion Design and Implementation.** Details the custom Python pipeline: its modular architecture, core `Sentence` and `Token` classes, and auditable logging system. It explains the two-phase process—(1) rule-based morphosyntactic mapping of POS tags and features, and (2) dependency restructuring—together with label mapping and post-processing that ensure full UD conformance.
- **Chapter 5: Validation and Outcomes.** Introduces the validation workflow based on the official UD validator, including the iterative procedure and validation dataset; reports the conformance results and treatment of remaining edge cases; and presents the final MPDT-UD 1.0 treebank, summarizing its size, data split, licensing, and integration into the UD ecosystem.
- **Chapter 6: Applications and Cognitive Science Perspective.** Positions the converter and MPDT-UD as reusable infrastructure, sketches applications in historical syntax and diachronic NLP, and relates the resource to questions about processing constraints and category change over time, concluding with proposals for future extensions of the pipeline and treebank.

Chapter 2

Background

This chapter provides the essential background for understanding the Middle Polish Dependency Treebank conversion to Universal Dependencies. It begins with the theoretical foundations of dependency grammar and its specific Polish manifestation in the Polish Dependency Bank (PDB) scheme (Section 2.1). Then it outlines Universal Dependencies as the target framework, highlighting its advantages for cross-linguistic research (Section 2.2). Finally, it describes the key resources: KorBa as the source corpus and MPDT as the dependency-annotated dataset that forms the input to our conversion pipeline (Section 2.3).

2.1. Dependency Grammar

Dependency grammar is a theory of syntactic structure organized around asymmetric governor–dependent relations. A *dependency* links two lexical items: a *governor* that selects and constrains a dependent, and a *dependent* that is licensed by the governor. One item can be a governor for multiple dependents, but each dependent has a single governor. Sentence structures are modeled as directed trees whose nodes correspond to tokens and whose edges encode these governor–dependent links. The tree has a single *root* (a node with no governor), and every other node is reachable from it along directed edges. In addition to purely structural links, dependency grammar is used here in a morphosyntactic sense, focusing on grammatical relations rather than semantic or prosodic dependency representations.

The dependency scheme used in Middle Polish follows the conventions established for the Polish Dependency Bank (PDB), which is adapted specifically for Polish syntax (Wróblewska 2023). The PDB tagset adapts the NKJP tagset (Przepiórkowski et al. 2012). The PDB annotation scheme uses a comprehensive set of part-of-speech categories and dependency relations designed specifically for Polish morphosyntax.

The PDB tagset includes the following part-of-speech categories:

- **Nouns:** `subst` (noun), `depr` (depreciative noun)
- **Pronouns:** `ppron12` (non-third person pronoun), `ppron3` (third person pronoun), `siebie` (reflexive pronoun)
- **Adjectives:** `adj` (adjective), `adja` (ad-adjectival adjective), `adjc` (predicative adjective), `adjp` (prepositional adjective)
- **Verb forms:** `fin` (finite non-past), `praet` (past tense), `imps` (impersonal), `impt` (imperative), `inf` (infinitive), `aglt` (agglutinate of ‘być’), `bedzie` (future form of ‘być’), `winię` (modal verbs like ‘winię’), `pred` (predicative), `ger` (gerund), `pcon` (contemporary adverbial participle), `pant` (anterior adverbial participle), `pact` (active adjectival participle), `ppas` (passive adjectival participle)
- **Numerals:** `num` (cardinal numeral), `numcomp` (numeral compound)
- **Conjunctions:** `comp` (subordinating conjunction), `conj` (coordinating conjunction)
- **Other categories:** `adv` (adverb), `brev` (abbreviation), `dig` (Arabic numeral), `romandig` (Roman numeral), `emo` (emoticon), `fill` (filler), `frag` (fragment), `interj` (interjection), `interp` (punctuation), `part` (particle), `prep` (preposition), `ign` (unrecognized form)

The PDB annotation scheme distinguishes several classes of dependency relations:

- **Core arguments:** `subj` (subject), `obj` (direct object), `obj_th` (thematic object), `comp` (complement), `comp_fin` (finite clause complement), `comp_inf` (open clause [*infinitive*] complement), `comp_ag` (agent complement)
- **Adjuncts and modifiers:** `adjunct` with semantic subtypes such as `adjunct_temp` (temporal), `adjunct_loc` (locative), `adjunct_dur` (duration), `adjunct_caus` (causal), `adjunct_mod` (manner), `adjunct_emph` (emphatic particle), `adjunct_compar` (comparative)
- **Predicate-related:** `pd` (predicative expression), `aux` (auxiliary), `neg` (negation), `refl` (reflexive)
- **Coordination:** `conjunct` (coordinated element), `pre_coord` (pre-coordinator)
- **Multiword expressions:** `mwe` (multiword expression), `ne` (named entity), `ne_foreign` (foreign named entity)

- **Special relations:** `punct` (punctuation), `vocative` (vocative), `orphan` (orphaned dependent), `discourse` (discourse marker), `parataxis` (parataxis), `aglt` (mobile inflection), `imp` (imperative marker), `cond` (conditional clitic), and `root` (sentence root)

The example dependency trees below illustrate the scheme of a PDB-annotated sentence alongside its UD counterpart, showing the structural differences.




Figure 1: Example dependency tree in the PDB format

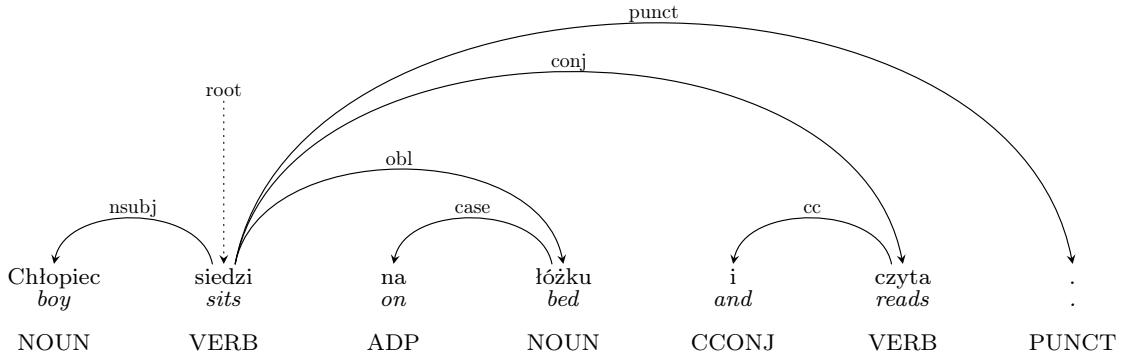


Figure 2: A dependency tree of the same sentence in UD format

Dependency formalisms differ on certain design choices (e.g., whether adpositions are heads or dependents inside adpositional phrases; how to encode coordination; whether and how to mark valency vs. modification). The PDB scheme takes specific positions on these issues, treating prepositions as heads (note the *on*→*bed* relation in Figure 1), using a coordination-centric approach where conjunctions govern coordinated elements (*i* being the `root` of both *siedzi* and *czyta* in Figure 1).

2.2. Universal Dependencies

Universal Dependencies (hereafter UD) is a cross-linguistic annotation framework designed to harmonize morphosyntactic and syntactic representations across languages

within a dependency-based, lexicalist model (Nivre et al. 2020; de Marneffe et al. 2021). UD serves as both a theoretical framework and a practical collection of treebanks—currently the largest repository of over 200 treebanks for more than 150 languages.¹ It is widely adopted in NLP and linguistic typology studies, and is maintained by an open community with regular releases.

Annotation scheme: The scheme provides three aligned layers for sentence-level annotation:

1. **Tokenization.** UD defines dependencies between *syntactic words*. To handle orthographic contractions or clitic clusters, it uses *multiword tokens*, ensuring a faithful word-level analysis. A multiword token is a single orthographic unit that is split into multiple syntactic words, each receiving its own morphological analysis and syntactic function.

For example, Middle Polish *kiedym* ('when I') is annotated as:

14-15	kiedym	-	-	...
14	kiedy	kiedy	ADV	...
15	m	być	AUX	...

Here, the single orthographic token *kiedym* (ID 14-15) splits into two syntactic words: *kiedy* 'when' (ID 14) and *m* (mobile inflection form of 'I am', ID 15).

Similarly, *jeszcześ* ('still you are') becomes:

7-8	jeszcześ	-	-	...
7	jeszcze	jeszcze	PART	...
8	ś	być	AUX	...

2. **Morphology.** Each syntactic word is associated with a LEMMA, a universal part-of-speech tag (hereafter part-of-speech tag=POS; universal part-of-speech tag=UPOS) from a fixed 17-tag set, and a bundle of FEATS (morphological features). The UPOS tags cover open-class words (adjectives ADJ, adverbs ADV, interjections INTJ, nouns NOUN, proper nouns PROPN, verbs VERB), closed-class words (adpositions ADP, auxiliary verbs AUX, coordinating conjunctions CCONJ, determiners DET, numerals NUM, pronouns PRON, particles PART, subordinating conjunctions SCONJ), and other categories (punctuation PUNCT, symbols SYM, other X). UD v2 standardized features and values across languages and clarified tag boundaries, e.g., extending auxiliary verbs to copulas and tense-aspect-mood particles while narrowing particles. The list of UPOS categories is available on the UD webpage.²

¹Universal Dependencies, <https://universaldependencies.org>, accessed 2025-10-10.

²Universal Dependencies POS tags: <https://universaldependencies.org/u/pos/index.html>

3. **Syntax.** The syntactic layer is a single-rooted tree with possible 37 universal dependency relations organized according to functional and structural categories. Sentence structures are modeled as directed trees according to the principles of dependency grammar as described in 2.1. Relations include:

- core arguments (nominal subject `nsubj`, direct object `obj`, indirect object `iobj`, clausal subject `csubj`, clausal complement `ccomp`, open clausal complement `xcomp`),
- non-core dependents (oblique `obl`, dislocated element `dislocated`, adverbial clause modifier `advcl`, adverbial modifier `advmod`, discourse element `discourse`, auxiliary `aux`, copula `cop`, vocative `vocative`, expletive `expl`, marker `mark`),
- nominal dependents (nominal modifier `nmod`, numeral modifier `nummod`, adjectival modifier `amod`, determiner `det`, case marker `case`, classifier `clf`, clausal modifier of noun `acl`, appositional modifier `appos`),
- coordination (conjunct `conj`, coordinating conjunction `cc`),
- multiword expressions (fixed `fixed`, flat `flat`),
- special relations (list element `list`, parataxis `parataxis`, orphan `orphan`, punct `punct`, root `root`, overridden disfluency `reparandum`, relation ‘goes with’ `goeswith`, other dependent `dep`).

The framework also allows language-specific subtypes (e.g., `nsubj:pass` for passive subjects, `det:poss` for possessive determiners) and defines semi-mandatory subtypes that should be used when the relevant phenomenon exists in the language. A full list of relations and subtypes, along with their descriptions, is available in the UD webpage.³

In addition to the *basic* representation, UD also defines an *enhanced* graph that adds extra arcs (and occasionally null nodes) to capture phenomena such as shared dependents in coordination, control and raising, relativization, and ellipsis. In Figure 2, the basic tree structure is shown; an enhanced representation would add an additional edge to represent the dependent (in this case: the subject) of *czyta* ('reads') as also being the *Chłopiec* ('boy'), as shown in figure Figure 3.

³Universal Dependencies relations list: <https://universaldependencies.org/u/dep/index.html>

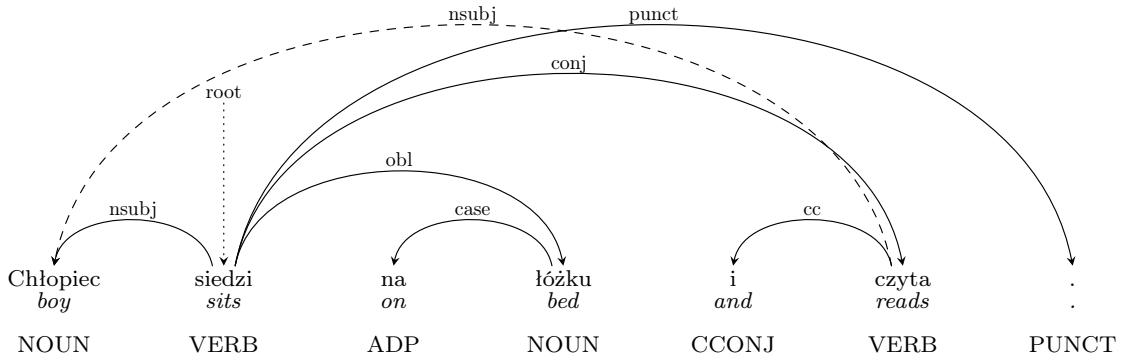


Figure 3: A dependency tree with enhanced dependencies (dashed lines)

Format: For practical implementation and data sharing, UD annotations must be encoded in a standardized format. UD uses the CoNLL-U format, a ten-column tabular specification with the fields:

- ID - a syntactic word index (or range for multiword tokens);
- FORM - the surface form;
- LEMMA - the dictionary form;
- UPOS - the universal POS tag;
- XPOS - a language-specific POS tag;
- FEATS - a pipe (|) separated list of morphological features;
- HEAD - the index of the head syntactic word (or 0 for the root);
- DEPREL - the dependency relation to the head;
- DEPS - for enhanced dependencies;
- MISC - for miscellaneous annotations.

Here is a CoNLL-U snippet for the sentence “Chłopiec siedzi na łóżku i czyta.”, with the enhanced dependencies.

# sent_id = test-sentence									
# text = Chłopiec siedzi na łóżku i czyta.									
1 Chłopiec	chłopiec	NOUN	subst	Gender=Masc Number=Sing Case=Nom		2	nsubj	-	-
2 siedzi	siedzieć	VERB	fin	Aspect=Imp Mood=Ind Tense=Pres Person=3 Number=Sing	0	root	-	-	-
3 na	na	ADP	prep	AdpType=Prep Case=Loc	4	case	-	-	-
4 łóżku	łóżko	NOUN	subst	Gender=Neut Number=Sing Case=Loc	2	obl	-	-	-
5 i	i	CCONJ	conj	-	2	cc	-	-	-
6 czyta	czytać	VERB	fin	Aspect=Imp Mood=Ind Tense=Pres Person=3 Number=Sing	2	conj	1:nsubj	-	-
7 .	.	PUNCT	interp	PunctType=Peri	2	punct	-	-	-

2.3. Middle Polish Linguistic Resources

2.3.1. KorBa

KorBa (Gruszczyński et al. 2022) – from Polish *Korpus Barokowy* ('Baroque Corpus') – is a 13.5-million-token corpus of Polish texts from 1601–1772, compiled from over seven hundred sources and annotated morphosyntactically (lemmas, POS, features). It is searchable via MTAS (Multi Tier Annotation Search; Brouwer et al. 2017), and provides parallel transliteration/transcription layers, structural and language markup, and rich metadata (period, region, text type, genre) that enable stratified analyses.

The corpus includes diverse text types ranging from literary works (epic poetry, drama, lyric poetry) to non-literary materials (scientific-didactic texts, persuasive writings, factual literature, official documents, press releases) and biblical texts. Geographically, texts span the Polish-Lithuanian Commonwealth, with approximately 27% of the corpus being of unknown origin. As shown in Figures 4 and 5, the corpus maintains careful balance across regions and text types to ensure representativeness of Middle Polish.

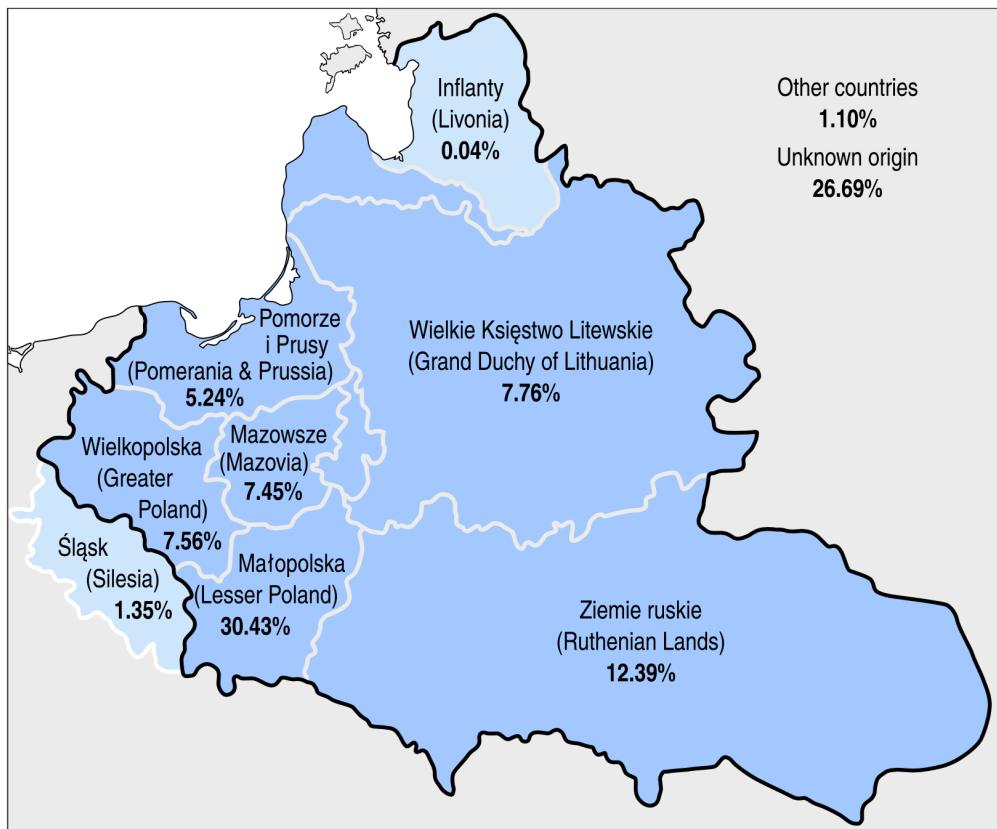


Figure 4: Geographical distribution of texts in the corpus displayed on the map of the Commonwealth after the Union of Lublin of 1569. Source: Gruszczyński et al. (2022), p. 315, CC BY 4.0.

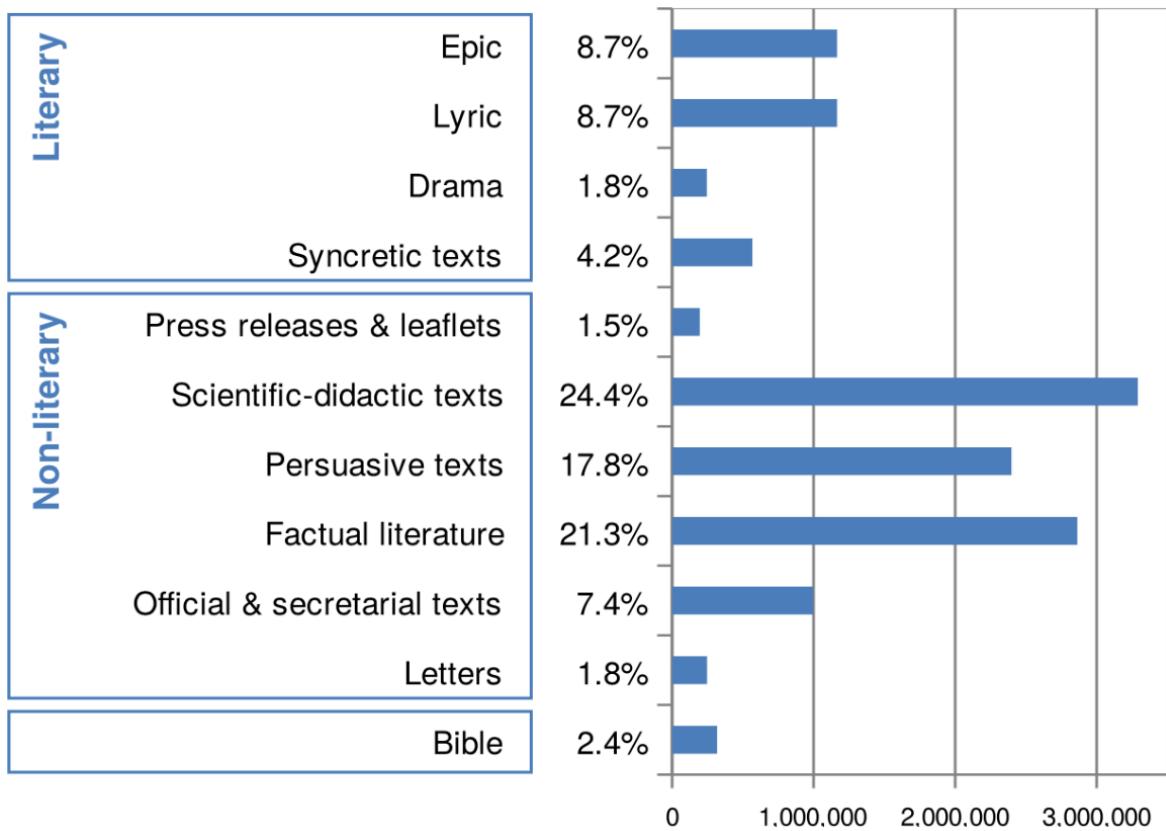


Figure 5: Types of texts in KorBa. Source: Gruszczyński et al. (2022), p. 316, CC BY 4.0.

2.3.2. MPDT

The Middle Polish Dependency Treebank (MPDT) is a manually curated, syntactically annotated subset of the KorBa corpus, capturing key syntactic phenomena of 17th–18th-century Polish texts. The sentences are from the manually annotated part of KorBa, whose careful pre-processing provides reliable morphosyntactic annotation and balanced coverage across genres and periods. In its current form, MPDT represents the first systematic attempt at syntactic annotation of Middle Polish and therefore in the current version excludes poetry and sentences with Latin insertions, while limiting sentence length to 10–50 tokens, with the average sentence length being 23 tokens (Wieczorek 2025).

The annotation workflow consists of the following steps:

1. **Automatic pre-annotation.** Two parsers trained on contemporary PDB data (MaltParser, COMBO) generate initial dependency analyses.
2. **Manual correction.** Two linguist annotators independently revise parser outputs, leveraging complementary error profiles.

3. **Adjudication.** Conflicting annotations are resolved by an adjudicator to produce a single gold-standard tree.
4. **Formatting.** Final annotations are encoded in CoNLL-X⁴ with KorBa’s extended tagset (e.g., dual number Dual).

Corpus statistics

- Total sentences: 2 018
- Total tokens: 47 273
- Distinct POS tags: 45
- Distinct dependency relations: 27
- Non-projective edges: 3 748 across 879 sentences
- Average sentence length: 23.43 tokens

Figure 6 presents the 20 most frequent MPDT POS tags, highlighting the prominence of nouns (**subst**: 11,374 occurrences), punctuation (**interp**: 7,971), adjectives (**adj**: 5,315), and prepositions (**prep**: 4,391).




Figure 6: Top 20 MPDT POS tag frequencies

⁴CoNLL-X is the predecessor of the CoNLL-U format (Buchholz and Marsi 2006)

Figure 7 shows the distribution of the top 20 dependency relation types. Adjuncts (`adjunct`: 13,276) and complements (`comp`: 8,539) are most common, followed by punctuation (`punct`: 6,896), coordination elements (`conjunct`: 6,071), and core arguments (`obj`: 3,423; `subj`: 2,286).

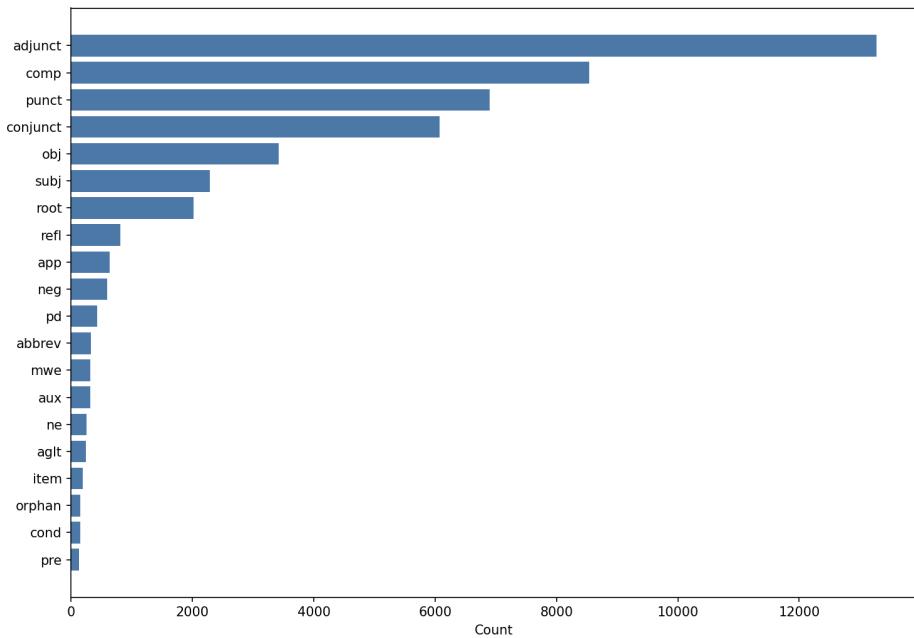


Figure 7: Top 20 MPDT dependency relation base frequencies

Chapter 3

Linguistic Features of Middle Polish

This chapter characterizes the linguistic system of Middle Polish as represented in the KorBa corpus and the Middle Polish Dependency Treebank (MPDT). It outlines key differences from modern Polish orthography and punctuation (Section 3.1), morphology (Section 3.2), and syntax (Section 3.3), emphasizing those that directly affect dependency annotation and conversion to Universal Dependencies (UD).

Note. Unless otherwise indicated, all Examples in this chapter and the chapters that follow (Chapters 3–6) are drawn from the Middle Polish Dependency Treebank (MPDT).

3.1. Orthography and Punctuation

3.1.1. Orthography and Transliteration

The KorBa corpus preserves two parallel orthographic layers: *transliteration* (a faithful rendering of the historical text) and *transcription* (a normalized spelling approximating contemporary Polish orthography). As noted in the KorBa manual, transliteration reflects the original graphic form of 17th–18th-century sources, while the transcription adapts them to modern conventions, keeping key phonetic and morphological features of Middle Polish.

Middle Polish orthography was far from standardized. The same word could appear in several spelling variants, sometimes even within a single text. Graphemes were often used interchangeably (e.g., *i/y*, *u/v*, *ć/ci*, *rz/ż*). Long vowels (*á, é*) and palatalization (*ć, ź, ś,ń*) were marked inconsistently. The KorBa transliteration layer preserves this variation, while the transcription layer normalizes it (e.g., *rodźicow* → *rodziców*).⁵

Orthographic conventions also influenced tokenization. Many expressions that are today written separately were then written together, and vice versa. For example:

⁵See Gruszczyński et al. (2022), p. 317.

- Historical joint writing *z chęci*, modern *z chęci* ('from willingness')
- Historical separate writing *dla tego*, modern *dlatego* ('because', lit. 'for this')

Following Wieczorek (2020), the MPDT treats historical spacing as evidence but bases syntactic analysis on function. Two frequent scenarios occur. (i) Historically fused sequences that correspond to ordinary prepositional phrases are split into their syntactic parts; the preposition and its nominal complement are annotated as in regular usage (see Example 1). (ii) Historically separate sequences that function as a single connective or adverbial are kept as two tokens but marked as a fixed multiword unit in MPDT; the unit then receives a single clausal function (see Example 2).

Example 1:

*Poniewoli musiała zbita na łóżu leżeć u rodziców/ ale **z chęci** obiecała sobie inną okazję/ iż jej od przedsięwzięcia nie oderwał.*
 ('Against her will she had to lie beaten on the bed at her parents'/ but of her own will she promised herself another occasion/ that they would not tear her away from her undertaking.')

In Example 1, *z chęci* is analyzed in MPDT as the preposition *z* with the genitive noun *chèci*: the preposition functions as a modifier of the non-finite predicate, and the noun is its complement.

Example 2:

*Tak jest: i **dla tego** tak się poniżył.*
 ('Indeed: and therefore he humbled himself thus.')

In Example 2, *dla tego* forms a lexicalized connective in MPDT: *dla* serves as the syntactic head of the unit that modifies the clause, and *tego* is linked to it as a multiword element within that unit.

3.1.2. Punctuation

As described by Wieczorek (2025), punctuation in Middle Polish reflected the rhythm and pauses of speech rather than syntactic boundaries. Marks were used inconsistently and sometimes idiosyncratically: slashes (/) often functioned as commas, semicolons as commas, and colons as semicolons or dashes. Conversely, long unpunctuated stretches also occur.

During syntactic annotation, punctuation is interpreted according to its syntactic function, not its original graphic mark. For instance, a slash (/) that introduces a new clause is annotated as *punct*.

Example 3:

*Powstawszy raz z bárzo ćięszkiew choroby/ ták rzekł Nie nagorzey śię
zemnq stálo: Bo mię chorobá vpomniálá/ ábym się w pyche nie podnošíł/
ponieważem iest śmiertelny.*

(‘Having once recovered from a very severe illness/ he said thus: It did not go too badly with me: For the illness reminded me/ that I should not lift myself up in pride/ since I am mortal.’)

As noted above, the slash in Example 3 functions as a clause delimiter and is therefore annotated as `punct`.

3.2. Morphology

The morphological system of Middle Polish differs significantly from the modern language, both in its inventory of forms and in category values. These distinctions were codified in the KorBa 2.0 tagset and later adopted in the MPDT.

3.2.1. Additional Parts of Speech and Morphological Features

The Middle Polish tagset introduces several categories and feature values that are rare or absent in contemporary Polish; below those are highlighted with direct impact on UD mapping.

(a) Short-form adjectives (adjb). These forms—e.g., *żyw* (‘alive’), *godzien* (‘worthy’)—are indeclinable or partially declined adjectives, often used predicatively without the copula. In UD they are mapped to `UPOS=ADJ` with `Variant=Short`.

Example 4:

Iak d ugo ia żyw iestem, żyie P an moy poty, Czui  bol y wesoło  , czui  y klopoty.

(‘As long as I am alive, my Lord lives likewise; I feel pain and joy, I feel troubles as well.’)

Example 5:

Chcesz si  zemnq r wna : nie godziene  tego.

(‘You want to match yourself with me: you are not worthy of this.’)

In modern Polish, the short form *żyw* from Example 4 would be considered archaic or poetic; the modern equivalent is *żyw*. The word *godzien* from Example 5 still exists, along with a few other, like *pewien* (‘certain’), however their usage is now limited, and the standard forms are *godny*, and *pewny*.

(b) Short passive participles (ppasb). Uninflected short passive participles—e.g., *zbawion* ('saved'), *pisan* ('written')—co-occur with finite forms of *być*. In UD they are annotated as UPOS=ADJ with VerbForm=Part, Voice=Pass, Variant=Short.

Example 6:

Kto vvierzy, á okrzci się, zbawion będącie, ále kto nie vvierzy będącie potępiion.

('Whoever believes and is baptized will be saved, but whoever does not believe will be condemned.')

Example 7:

Pisan na zamku pileckim, dnia 23 miesiąca lipca, roku Pańskiego 1620.

('Written at the castle of Pilec, on the 23rd day of July, in the Year of Our Lord 1620.')

In modern Polish, short forms from Examples 6 and 7 are archaic; the standard forms are fully inflected *zbawiony*, *potępiony*, *pisany*.

(c) Past participles (ppraet). Forms such as *osłabiałe* ('weakened'), *opuchłyymi* ('swollen'), *zasiniałyymi* ('bruised/blue-tinged') represent an older stage of adjectival participles derived from past tenses, intermediate between ppas and pact. In UD they are mapped to UPOS=ADJ with VerbForm=Part and Voice=Pass.

Example 8:

Częstokroć abowiem były widane z twarzami opuchłyymi/ zásiniąłyymi.

('For often they were seen with swollen/ bruised faces.')

In modern Polish, the past participle forms are still in use, but some are archaic or poetic. Looking at words from Example 8 *opuchłyymi* would be rather replaced by *opuchniętymi*, while *zasiniałyymi* is still acceptable.

(d) Dual number (du). Middle Polish still preserved dual forms for certain nouns, numerals, adjectives, and verbs. The KorBa manual documents the explicit tag du. These forms gradually merged with the plural after ca. 1740, though fossilized duals like *ręce*, *oczy* survive in modern Polish (singular *oko* ('eye') → plural *oczy* when referring to the organ, but also pl. *oka* when used in other sense, e.g., *oka w rosole* ('eyes in the broth'); similarly singular *uchو* ('ear') → plural *uszy*, when about body parts, or pl. *uchy*, when referring to cup handles). In UD, dual forms are annotated with Number=Dual.

Example 9:

6. Po przepędzonych przez **dwie lecie** tych okrutnych bolesciach, pokazał się iey Pan mowiąc: *Iż bez lat pięć nie miałaby iadać ani mięsa. ani nabiału.*
 ('6. After two years spent in these cruel pains, the Lord appeared to her saying: That for five years she should not eat either meat or dairy.'')

In Example 9, *dwie lecie* is dual accusative of *dwa* (two) and *lato* ('summer; year'). In modern Polish, the dual form is archaic; the standard form (both the nominative and accusative) is *dwa lata*.

3.2.2. Gender System and Declension

The masculine gender system in Middle Polish was less differentiated than in the modern language. KorBa distinguishes three values: **m** (general masculine), **manim1** (masculine personal), and **manim2** (masculine non-personal). In early texts, these values overlap; many forms do not yet reflect consistent distinctions in case endings. For example, *ptaki* and *ptacy* ('birds') alternate depending on context.

Example 10:

6. Vbogáciłeś ich chybkością i lotem nád wszystkie loty przedzym i bystrzejszym/ i bystрым ták/ iż i strzały/ i **ptaki**/ i pioruny poprzedzác/ á wszystkie rzeczy/ mury/ skály/ przenikać mogą.

('You have enriched them with speed and with flight swifter and sharper than all flights/ so that even arrows/ and birds/ and thunder they can outpace/ and penetrate all things/ walls/ rocks.'')

Example 11:

122. Czemu **ptacy** ktorzy ogona nie máją długie nogi mają?
- ('Why do birds that do not have a tail have long legs?')

Example 10 illustrates the use of *ptaki* in the general masculine category (**m**), while Example 11 uses the masculine personal value *ptacy* (**manim1**).

3.3. Syntax

3.3.1. Word Order and Non-projectivity

Middle Polish syntax exhibits high flexibility of word order, frequent inversion, and long-distance dependencies. As noted by Wieczorek (2025), discontinuous structures—especially in noun phrases with adjectival modifiers—often yield non-projective

trees. The contrast between a linear and a discontinuous configuration is illustrated in Figures 8 and 9.




Figure 8: Linear order (no crossing edges)

Source: adapted from Wieczorek (2025), Fig. 6, p. 12.




Figure 9: Discontinuous order with crossing edges between *dobra* and *fundacyja*

Source: adapted from Wieczorek (2025), Fig. 7, p. 12.

These inversions complicate automatic parsing and were one challenge for explicit rule-based conversion to UD.

3.3.2. Predicate Ellipsis

As noted by Wieczorek (2025), it is rare in modern Polish for sentences to lack a predicate (at least in texts written in careful language), but this was quite common in 17th–18th-century Polish. In dependency analysis, the predicate is considered the centre of the sentence (`root`)—most often a finite verb. In the absence of a predicate, another sentence element serves as the centre. Most often, this centre becomes the subject, which receives the label `root` instead of `subj`, as in Figure 10.




Figure 10: Predicate ellipsis: nominal head as **root**

Source: adapted from Wieczorek (2025), Fig. 8, p. 13.

3.3.3. Clause Linking and Subordination

Middle Polish frequently employs conjunctions that have since changed meaning, an example of which could be the token *jako*. It is frequently employed in two functions: (i) as a comparative/similative marker in the sense of modern *jak* ('like/as; when/how'), and (ii) in the modern-like role/identity sense 'as'. The readings are illustrated in Examples 12–14.

Example 12:

Ale jako nowi Obywatele tam przybywać poczeli z Kir, czy Syr Kraiu w Medii leżącego, Kirya, to Syria zwać się poczęła.
 ('But as/when new inhabitants began to arrive there from the land of Kir, or Syr, lying in Media, Kirya then began to be called Syria.')

Example 13:

Iako Roża rozpuszcza z przyrodnego swego zapach przyjemny, tak Serce dobrotzynne wydaie bez przyniewolenia uczynki dobre.
 ('As a rose by its nature gives off a pleasant fragrance, so a charitable heart produces good deeds without compulsion.')

Example 14:

Ja zaś w tych terminach stawam jako mediator, prowadząc do zgody obiekie strony.
 ('And I, for my part, in these proceedings stand as a mediator, leading both sides to agreement.')

In Example 12, *jako* functions as a clause linker with temporal meaning ('when'). In Example 13 it introduces a comparative clause ('how/as'), and in Example 14 it introduces a role/identity complement ('as (in the role of)'). In modern Polish, the first two functions are typically expressed with *jak*, while the third remains *jako*.

3.4. Summary

Middle Polish exhibits substantial divergence from modern Polish in orthography, morphology, and syntax:

- Orthography: inconsistent, variable, often merging or splitting tokens differently from modern norms.
- Punctuation: prosodic rather than syntactic, with slashes and colons used irregularly.
- Morphology: additional forms (**adjb**, **ppasb**, **ppraet**), productive dual number, and fluid gender distinctions.
- Syntax: high non-projectivity, frequent inversion, ellipsis, and loose coordination.

These properties directly inform the design of the MPDT → MPDT-UD conversion pipeline, motivating special conversion rules and additional validation layers to preserve linguistic authenticity while ensuring formal compatibility with Universal Dependencies.

Chapter 4

Conversion Design and Implementation

This chapter details the design and implementation of the MPDT→MPDT-UD conversion pipeline. The conversion is a complex, multi-stage process, divided into two primary phases: 1. morphosyntactic mapping (Section 4.3) and 2. dependency tree transformation (Section 4.4). Before detailing them, the chapter first describes the overall architecture of the converter, including its custom data structures and the code repository layout (Section 4.1), the high-level processing workflow (Section 4.2), and at the end—the auditable logging system that fulfills research goal **(R2)** (Section 4.5).

4.1. Converter Architecture and Environment

The entire conversion process is implemented in Python, leveraging a custom-built environment designed for traceability and modularity. The code is publicly available in a GitHub repository: <https://github.com/kvmilos/MPDT-to-UD-converter>.

4.1.1. Core Data Structures

The environment is built around core data structures, `Sentence` and `Token` classes, defined in `utils/classes.py`. A key design choice is that each `Token` object stores both the original MPDT annotation and the new, converted UD annotation in parallel. This allows conversion rules to access the original, unmodified MPDT context at any stage, which is crucial for resolving ambiguity during the complex dependency transformation phase.

The `Token` class is also equipped with numerous helper properties and methods to simplify the writing of conversion rules, such as methods accessing the governor in the new UD tree if it is present, and accessing the old one otherwise, or traversing the tree to find specific dependents or governors via certain relations.

In parallel, the `Sentence` class provides the structural container for all `Token` objects belonging to a single sentence. Beyond simple storage, it offers functionality that is essential for conversion: it maintains a fast token lookup table (`dict_by_id`), stores sentence-level metadata (e.g., `# sent_id`, `# text`), and exposes utility methods for navigating the dependency tree. These include retrieving the root token, or iterating over tokens in different orders. By centralizing this behaviour in a dedicated class, the converter ensures that all modules operate over a consistent and fully accessible representation of sentence structure.

4.1.2. Project Repository Structure

The converter's code is organized into modules based on functionality. The main modules handle the top-level pipeline, morphosyntactic conversion, dependency conversion, and utilities. The structure of the repository is as follows (directories are shown in blue):

```

ud_converter/
    └── converter.py
    └── morphosyntax/
        └── pos_categories/
            ├── preconversion.py
            ├── conversion.py
            └── postconversion.py
        └── morphosyntax.py
    └── dependency/
        └── structures/
            ├── labels.py
            ├── edges.py
            ├── preconversion.py
            ├── conversion.py
            └── postconversion.py
    └── utils/
        ├── classes.py
        ├── constants.py
        ├── io.py
        └── logger.py
    └── data/
    └── logs/

```

The main components are:

- `converter.py`: The main executable script that orchestrates the entire conversion pipeline.
- `morphosyntax`: The package for Phase 1 (morphosyntactic conversion).

- `pos_categories`: Contains a separate module for most MPDT XPOS tags (e.g., `subst.py`, `adj.py`) to handle its specific conversion rules.
- `dependency`: The package for Phase 2 (dependency conversion).
 - `structures`: Contains modules for restructuring specific syntactic constructions (e.g., `coordination.py`, `prepositional.py`).
- `utils`: Utility package with helper modules for the entire application.
 - `classes.py`: Defines the core `Sentence` and `Token` objects.
 - `constants.py`: A central store for all static mappings (e.g., features, lemmas).
 - `io.py`: Handles reading input `.conll` and `.json` files and writing the output `.conllu` file.
 - `logger.py`: Implements the auditable logging system, including the `ChangeCollector` and `LoggingDict` classes.
- `data/`: Default directory for input and output data files.
- `logs/`: Default directory where the detailed conversion logs are saved.

4.2. Conversion Pipeline

The converter is designed as a sequential pipeline, executed from the main `converter.py` script. From the user's perspective, the process consists of four main stages:

- (1) **Data Loading:** The pipeline begins by reading two input files: the MPDT treebank in its `.conll` format and a corresponding `.json` metadata file. The data is loaded into the custom `Sentence` and `Token` objects.
- (2) **Phase 1: Morphosyntactic Conversion:** The first processing stage performs a rule-based conversion of the MPDT morphosyntactic annotations into their UD counterparts.
- (3) **Phase 2: Dependency Conversion:** The second processing stage transforms the syntactic structure of the trees. This highly contextual phase converts the MPDT dependency relations to UD relations and restructures the tree topology to conform to UD guidelines.
- (4) **Output Generation:** Finally, the converted `Sentence` and `Token` objects, now populated with UD annotations, are written to a single output `.conllu` file.

4.3. Phase 1: Morphosyntactic Conversion

The first processing phase, handled by the `morphosyntax` module, converts the MPDT `XPOS` tags and morphological features into their `UPOS` and `FEATS` counterparts. This phase is executed as a three-step sub-pipeline for each sentence.

4.3.1. Pre-conversion

First, a set of lemma-based rules are applied to handle specific lexical items whose categorization overrides the more general `XPOS`-based rules. For example:

- Conjunctions like *niz* ('than'), *jakby* ('as if'), and *niczym* ('like'), which introduce comparisons, are unambiguously mapped to `SCONJ` (subordinating conjunction) and assigned the feature `ConjType=Comp` to mark this comparative function.
- The lemma *temu*, when used as a postposition (e.g., *dwa lata temu* 'two years ago'), is mapped to `ADP` (adposition) and assigned the feature `AdpType=Post` to explicitly mark it as a postposition, distinguishing it from the standard prepositional form.
- Words with an initial capital letter that are not otherwise classified (e.g., as verbs) are provisionally tagged `PROPN` (proper noun). This rule helps correct cases where a proper noun was ambiguously tagged as a common noun (`subst`).

4.3.2. Core POS Conversion

Next, the main conversion logic maps the MPDT `XPOS` tag of each token to its corresponding `UPOS` tag and `FEATS`. The converter dispatches each token to a dedicated function based on its `XPOS` tag.

This design handles both simple and complex conversions. For instance, while `conj` (coordinating conjunction) almost always becomes `CCONJ`, the `subst` (noun) tag requires more logic: most `subst` tokens are mapped to `NOUN`, but the converter first checks for pronominal lemmas (e.g., *kto*, *co*, *nikt*) and maps these to `PRON` (pronoun) with the appropriate `PronType` feature.

This module also handles the specific Middle Polish phenomena described in Chapter 3:

- `adjb` (short adjective) is mapped to `UPOS=ADJ` and given the feature `Variant=Short`.
- `ppasb` (short passive participle) is mapped to `UPOS=ADJ` with the features `VerbForm=Part`, `Voice=Pass`, and `Variant=Short`.
- The MPDT gender system is correctly mapped to the UD features (e.g., `manim1` to `Gender=Masc` and `Animacy=Hum`).

- The Middle Polish Number=Dual feature is preserved, as it is a valid feature in Universal Dependencies, even if absent in modern Polish.

4.3.3. Post-conversion

Finally, a sentence-level cleanup function performs two crucial tasks that require the original, non-tokenized text from the metadata:

1. **Reconstructing Multiword Tokens:** This step correctly formats clitic constructions that were already split into syntactic words in the input data. For example, for the Middle Polish word *kiedym* ('when I'), the input .conll file contains two separate token lines (*kiedy* and *m*). This function reads the original text, sees they are not space-separated, and inserts the required multiword token entry (e.g., 14-15 *kiedym* ...) before the syntactic words it spans, as shown in the example in Chapter 2.2.
2. **Annotating Spaces:** The same function analyzes the original text to add SpaceAfter=No to the MISC column for any token that is immediately followed by another token or punctuation mark without an intervening space. This is a requirement for the CoNLL-U format since Universal Dependencies v. 2.0.

4.4. Phase 2: Dependency Conversion

The second phase, managed by the dependency module, is significantly more complex. Unlike morphosyntax, dependency conversion is not token-local; rules must consider a token's governor, its dependents, and its siblings, often operating on the original MPDT structure, the partially converted UD structure, or both.

Many of the structural transformations were adapted from the principles established for the conversion of the contemporary Polish Dependency Bank (Wróblewska 2018; Wróblewska 2020), but were re-implemented to fit the custom pipeline and handle Middle Polish phenomena. The conversion follows a strict pipeline of restructuring, label mapping, and post-processing.

4.4.1. Structural Restructuring

The first and most critical step is to change the topology of the dependency tree to conform to UD principles. The converter applies a series of modules to handle specific syntactic constructions. The most fundamental transformations are:

- **Prepositional Phrases:** In MPDT, a preposition (prep) governs its nominal complement (comp). This structure gets inverted: the nominal complement be-

comes the head of the phrase. This new head inherits the original syntactic function from the preposition (e.g., the `adjunct_locat` relation from *na* in Figure 11 becomes the dependency for *koniu*). This relation is then mapped to a UD relation (e.g., `obl`). Finally, the preposition is re-attached to the noun with the `case` relation. This transformation is illustrated in Figure 11.




Figure 11: MPDT analysis (arcs above) and converted UD analysis (arcs below) for the fragment *siedzący na koniu rozpędzonym* ('sitting on a galloping horse').

- **Numeral Phrases:** Numeral expressions that govern their nouns in MPDT (e.g., as a `comp`) are restructured. In UD, the noun is promoted to be the head, and the numeral is re-attached as its dependent with the `nummod` relation. The noun phrase then attaches to the verb (e.g., as `obl`), as shown in Figure 12.




Figure 12: MPDT analysis (arcs above) and converted UD analysis (arcs below) for the clause *Dwa dni wisiał* ('He hung for two days').

- **Predicative expressions:** In MPDT, the copula (e.g., *być* 'to be' or the `pred` token *to*) is the head of the clause, governing the subject (with `subj`) and the non-verbal predicate (with `pd`). To comply with UD, the non-verbal predicate is promoted to be the head. The subject and the copular verb are then re-attached as dependents of this new nominal or adjectival head, receiving the UD relations `nsubj` and `cop`, respectively. This is illustrated in Figure 13.




Figure 13: MPDT analysis (arcs above) and converted UD analysis (arcs below) for the fragment *Główą tego królestwá iest Insułá* ('The head of this kingdom is Insułá').

- **Subordinate Clauses:** In MPDT, a subordinating conjunction (`comp`) often governs the predicate of its clause (e.g., as `comp_fin`). The converter inverts this, promoting the subordinate clause's predicate to be the head (which then attaches to the main clause predicate, or becomes the root), and demotes the conjunction to be a dependent of its clause's predicate with the `mark` relation. In Figure 14, this corresponds to replacing the MPDT chain $rzekł \xrightarrow{\text{comp}} že \xrightarrow{\text{comp_fin}} było$ with the UD configuration $rzekł \xrightarrow{\text{ccomp}} lekko$ and $lekko \xrightarrow{\text{mark}} že$.




Figure 14: MPDT analysis (arcs above) and converted UD analysis (arcs below) for the fragment *rzekł, że mu nigdy tak lekko nie było* ('he said that it had never been so easy for him').

- **Coordination:** In MPDT, the coordinating conjunction (`conj`) is the head of the coordinated elements (`conjunct`). This structure is rebuilt by promoting the *first* conjunct to be the head of the entire coordination. Subsequent conjuncts are attached to this first conjunct with the `conj` relation. The conjunction itself is re-attached to its *following* conjunct with the `cc` relation. In the upper tree of Figure 15, the coordinator *y* governs both conjuncts *odwagi* and *sercá*, each of which has its own modifier. In the lower tree, the first conjunct *odwagi* becomes the syntactic head of the coordination, the second conjunct *sercá* attaches to it as `conj`, and the conjunction *y* becomes a dependent of *sercá* with relation `cc`.




Figure 15: MPDT analysis (arcs above) and converted UD analysis (arcs below) for the fragment *odwagi swoiey y sercá rycerskiego* ('of one's courage and knightly heart').

4.4.2. Label Mapping

After the tree structure is finalized, a dedicated module traverses the tree and assigns a final UDEPREL to each token. This mapping is highly context-sensitive, using the UPOS of both the token and its new governor, as well as its original MPDT DEPREL. For example, the generic MPDT **adjunct** relation is mapped to a variety of UD relations:

- A nominal modifier of a noun (**adjunct** on a NOUN dependent) → **nmod**
- An adjectival modifier of a noun (**adjunct** on an ADJ dependent) → **amod**
- An adverbial modifier of a verb (**adjunct** on an ADV dependent) → **advmod**
- A prepositional phrase modifying a verb (**adjunct** inherited by a NOUN from a prep) → **obl**
- A clausal modifier of a verb (**adjunct** on a VERB dependent) → **advcl**

4.4.3. Correction and Post-processing

Finally, a series of cleanup scripts are run. One module ensures UD validation compliance by removing disallowed dependents (e.g., a **case** token cannot have its own dependents, so any punctuation attached to it is moved to its head).

Another module handles final tasks, such as disambiguating pronouns that are ambiguous between interrogative and relative (e.g., **PronType=Int,Rel** → **PronType=Int** or **PronType=Rel**) based on their new syntactic context. Most importantly for downstream use, this module generates the enhanced dependency graph (DEPS column) by propagating shared dependents in coordination. This contributes to research goal **(R3)** by improving evaluability and supporting enhanced UD representations.

4.5. Audibility and Processing Workflow

Beyond the core conversion logic, two key features of the converter are its auditable design and its straightforward user workflow.

4.5.1. Logging and Traceability

A core design principle of the converter is audibility, fulfilling research goal (**R2**). This is implemented via a custom logging system built into the `utils/logger.py` module.

A central `ChangeCollector` class gathers change events from all modules. To automate this, the core `Token.data` dictionary is implemented as a `LoggingDict`, a dictionary subclass that automatically calls `ChangeCollector.record()` whenever a value is set or changed.

Each log entry records the sentence ID, token ID, the specific module and function that triggered the change, and a message detailing the transformation (e.g., `upos changed from VERB to AUX`). This fine-grained logging (Contribution **C1**) proved invaluable for debugging, as it allows for a step-by-step reconstruction of how a token was processed and which rules fired. It was particularly critical for identifying and resolving rule conflicts during the complex dependency conversion phase.

4.5.2. Processing Workflow

From a user's perspective, the pipeline is executed via a single command. The converter takes the MPDT `.conll` file and the corresponding metadata `.json` file as input.

```
python converter.py input_file.conll output_file.conllu meta_file.json
```

The script processes each sentence and saves the result in the specified `output_file.conllu` in the valid CoNLL-U format, ready for validation and downstream use. The converter can be run in two modes. By default, it executes the complete pipeline, performing both phases of the conversion (morphosyntax and dependencies). However, if the user provides the `-tags-only` command-line flag, the pipeline omits Phase 2. This allows the user to generate a file with only the morphosyntactic conversion applied, leaving the original MPDT dependency structure intact.

4.6. Summary

This chapter has defined and implemented the core of the MPDT → MPDT-UD conversion pipeline:

- It specified a UD-oriented conversion strategy for MPDT, including the main structural transformations (prepositional phrases, numeral phrases, predicative expressions, subordination, and coordination), directly addressing research goal (**R1**).
- It described a modular, auditable implementation in Python, centred on custom `Sentence` and `Token` classes and an integrated logging system, thereby realizing research goal (**R2**) and delivering Contribution (**C1**) (the converter itself).
- It prepared the ground for formal UD conformance and evaluability (**R3**) by enforcing UD-compliant structures, mapping labels to the UD inventory, and generating enhanced dependencies; the quantitative conformance results are presented in Chapter 5.

Together, these elements produce a reusable conversion pipeline and a traceable implementation that underpins the validated MPDT-UD treebank discussed in the following chapter.

Chapter 5

Validation and Outcomes

This chapter presents the final, tangible outcome of the thesis: the initial public release of the **Middle Polish Dependency Treebank in Universal Dependencies format (MPDT-UD)**. Beyond reporting that the treebank passes the official validator, the chapter explains *how* validation shaped the converter’s design: validator feedback drove targeted, logged corrections rather than ad-hoc edits, yielding a reproducible and auditable pipeline.

The chapter is structured as follows. Section 5.1 introduces the validation workflow, including the iterative procedure, the validation dataset, and the official UD validator. Section 5.2 reports the conformance results and outlines how remaining edge cases were handled. Section 5.3 presents the final MPDT-UD 1.0 treebank, summarizing its size, data split, licensing, and integration into the UD ecosystem.

5.1. Validation Workflow and Formal Conformance

To fulfill research goal (**R3**)—ensuring that the converted treebank is formally correct and compatible with standard UD tools—a strict validation workflow was adopted. The entire MPDT-UD corpus was checked using the **official Universal Dependencies validator script validate.py**.⁶ The procedure was implemented as an iterative refinement cycle, applied to a fixed validation dataset and grounded in the behaviour of the official validator.

5.1.1. Iterative Validation Procedure

The validation workflow was integrated into the development process as an iterative, data-driven refinement cycle that directly influenced the converter’s architecture. Achieving the final zero-error conformance required repeatedly executing the following

⁶The validator is part of the UD tools repository; available at <https://github.com/UniversalDependencies/tools/blob/master/validate.py>

steps:

1. Running the full MPDT → MPDT-UD conversion pipeline.
2. Executing the `validate.py` script on the resulting `.conllu` file.
3. Collecting and classifying validator messages to identify systematic error types.
4. Implementing targeted adjustments in the relevant module (typically postconversion or label mapping) and repeating the cycle.

Throughout development, validator feedback served as an objective convergence criterion: changes to the pipeline were accepted only if they reduced violations without introducing new ones, and every change was captured by the converter’s logging facilities for later audit.

5.1.2. Validation Dataset

The validation was performed on the entire output of the conversion pipeline, i.e. the full MPDT available at the time of writing, converted to UD as MPDT-UD. All mentions in this chapter refer to this complete set.

Concretely, the validation dataset corresponds to the initial UD release of UD_Polish-MPDT (version 2.17). It contains **2,018 sentences**. The material is drawn from the manually annotated portion of the KorBa corpus (as described in Section 2.3) and thus inherits its genre balance and time-span.

Since Middle Polish is not a separate language branch within the UD project, the treebank was validated under the tagset and feature inventory of modern Polish (p1). This approach guarantees maximal compatibility with existing UD resources and ensures that the resulting data can be immediately processed by any standard UD-compliant parser or visualization tool.

5.1.3. Official UD Validator

The `validate.py` script performs a comprehensive, multi-layer verification of every sentence in a CoNLL-U file. It is implemented in Python and relies on the official UD feature inventories and relation lists, performing checks in three main categories:

- **Format compliance:** ensures that each sentence adheres to the CoNLL-U specification—exactly ten columns per token, valid multiword token ranges, correct comments, and consistent `SpaceAfter=No` annotation.
- **Morphological validity:** verifies that every UPOS tag, XPOS tag, and FEATS combination is legal according to UD v2.17 inventories (e.g., that `Aspect` features only

occur on verbal categories, and that `PronType` values correspond to pronouns or determiners).

- **Syntactic structure:** checks that the dependency tree is single-rooted, acyclic, and projective when required; that function words (e.g., tokens with relations `case`, `cc`, `mark`) do not have their own dependents; and that all `HEAD` indices refer to valid token IDs.

The validator also provides detailed diagnostic messages for every detected issue, grouped by error type and line number. This makes it suitable not only for final compliance testing but also for iterative debugging during development.

The script must be run with a language code (e.g., `-lang pl`) to check against treebank-specific feature and relation lists. It requires Python 3 and the `regex` and `udapi` modules. A typical invocation looked like:

```
python validate.py --lang pl MPDT-UD.conllu
```

5.2. Conformance Results

The final converted MPDT-UD treebank passes the official UD validator (`validate.py`) with **zero errors**. This 100% conformance fulfills the principal technical objective of the thesis (**R3**) and confirms that the resource is formally compatible with the Universal Dependencies standard.

This outcome is the result of the iterative, data-driven refinement cycle described in Section 5.1.1. Validator feedback was treated as an objective convergence criterion: pipeline changes were accepted only when they reduced the number of violations without introducing new ones, and all modifications were recorded by the converter’s logging facilities for later audit.

The emphasis throughout was on *principle-driven* cleanups aligned with UD guidance rather than ad hoc edits—for example, enforcing that function words do not bear dependents, attaching punctuation to appropriate content heads, and canonicalizing order-sensitive multiword relations. In this way, the final state of the converter encodes the corrections required to reach full conformance, rather than relying on one-off manual adjustments to the released data.

5.2.1. Handling Idiosyncratic Edge Cases

Most validation issues were solvable with generalized, context-sensitive rules that can be applied uniformly across the corpus. A small residual set of idiosyncratic cases—rare constructions not easily captured by broad heuristics—were addressed pragmatically to achieve full conformance on the released dataset. These interventions are minimal,

logged, and isolated, preserving reproducibility without overfitting the pipeline to particular sentences.

Conversion decisions for **POS** and **FEATS** are documented systematically in the accompanying specification.⁷ This document provides a stable reference for future extensions of the pipeline and for users who wish to understand or reuse the morphosyntactic mapping outside the present project.

5.3. Outcomes: the MPDT-UD 1.0 Treebank

The validated conversion pipeline yields the first public release of the Middle Polish Dependency Treebank in Universal Dependencies format (**MPDT-UD 1.0**). This section summarizes the size and internal structure of the treebank and briefly documents its integration into the Universal Dependencies ecosystem, including licensing and availability.

5.3.1. Corpus Size and Data Split

The MPDT-UD 1.0 treebank contains **2,018 sentences**, **46,670 surface tokens**, and **47,273 syntactic words**. The difference between tokens and syntactic words reflects the use of multiword tokens in the CoNLL-U representation. The corpus includes **564 multiword tokens** (across **359** distinct types), with an average of 2.07 syntactic words per multiword token. In addition, **8,217 tokens** (18%) are not followed by a space, which is a direct consequence of historical orthographic conventions such as clitic fusion and non-standard punctuation spacing.⁸

In line with the UD guidelines for treebanks in the 20K–110K word range, approximately 10K words were allocated to the test set, about 10% of the remaining data to the development set, and the rest to the training set. Sentences were randomly shuffled (using seed 42) and then assigned to the three subsets based on token-count quotas. The resulting split is:

- **Training set:** 1,433 sentences, 33,520 tokens.
- **Development set:** 162 sentences, 3,748 tokens.
- **Test set:** 423 sentences, 10,005 tokens.

The treebank does not preserve document boundaries or genre labels at the sentence level; sentences have been shuffled and redistributed purely for the purpose of forming these three subsets. At the corpus level, the metadata assign four broad genres to the resource as a whole: nonfiction, bible, legal, and fiction.

⁷See <https://github.com/kvmilos/MPDT-to-UD-converter/blob/main/MPDT.md>

⁸Summary statistics and tokenization details are publicly available via the UD_Polish-MPDT treebank page: https://universaldependencies.org/treebanks/pl_mpdt/index.html.

5.3.2. Release, Licensing, and UD Integration

MPDT-UD 1.0 is distributed within the Universal Dependencies project under the repository name **UD_Polish-MPDT**. It was first released as part of the **UD v2.17** release. The data are available both through the UD GitHub organization and via the standard UD release packages, alongside other Polish resources.

The treebank is released under the **Creative Commons Attribution–ShareAlike 4.0 (CC BY-SA 4.0)** license.⁹ In practical terms:

- **You may** share and adapt the data, including for commercial purposes.
- **You must** give appropriate credit, provide a link to the license, and indicate if changes were made.
- **ShareAlike:** if you remix, transform, or build upon the material, you must distribute your contributions under the same license as the original.
- **No additional restrictions:** you may not apply legal terms or technological measures that legally restrict others from doing anything the license permits.

These conditions align with the open-data practices of the UD project and make MPDT-UD 1.0 immediately reusable for historical syntax research, diachronic NLP experiments, and cross-linguistic studies that rely on UD-conformant training and evaluation data.

⁹License text and a plain-language summary are available at <https://creativecommons.org/licenses/by-sa/4.0/>.

Chapter 6

Applications and Cognitive Science Perspective

This chapter situates the MPDT→MPDT-UD converter and the resulting MPDT-UD treebank in a broader research context. It first discusses who can benefit from the tool and how it is packaged and licensed, then shows how it integrates with the Universal Dependencies ecosystem as part of a reproducible workflow (Section 6.1). It then sketches use cases in historical linguistics and diachronic NLP (Section 6.2), and finally connects the project to questions about processing constraints and category change in cognitive science (Section 6.3), concluding with directions for future work (Section 6.4).

6.1. Usefulness and Audience

The MPDT→MPDT-UD converter is the main software artifact produced in this project. While Chapter 5 focused on the resulting treebank (MPDT-UD 1.0) as a data resource, this chapter considers the converter itself as a reusable tool. The converter makes it possible to regenerate the UD version of the treebank from its MPDT source, to extend the coverage as MPDT grows, and to adapt the conversion strategy in a transparent, documented way. Because the pipeline is rule-based, modular, and fully logged, it is intended to be understandable and modifiable by other researchers working with historical Polish or related resources.

The following sections outline who can benefit from the converter, how it is packaged and licensed, and how it fits into a reproducible workflow for generating UD-conformant treebanks.

6.1.1. Who Benefits and How

The primary intended users of the converter are researchers and developers who work with Middle Polish data or similar historical resources:

- **MPDT and KorBa maintainers.** For the team curating MPDT and the underlying KorBa corpus, the converter provides a repeatable way of producing an up-to-date UD version whenever the source annotation is expanded or revised. Rather than editing the UD data directly, changes can be made in the MPDT layer and propagated through the pipeline, keeping the two representations in sync.
- **Historical treebank builders.** Other projects that annotate historical varieties of Polish or related Slavic languages can use the converter as a concrete template. Even if their annotation schemes differ, the basic separation between morphosyntactic mapping and dependency restructuring, combined with explicit logging, illustrates one way of organizing a conversion pipeline from a language-specific scheme into UD.
- **NLP practitioners and tool developers.** For parser developers and NLP practitioners interested in diachronic data, the converter makes it possible to regenerate training and evaluation sets under slightly different mapping choices (for example, alternative treatments of coordination or specific clause-linking strategies). This enables controlled experiments on how particular conversion decisions affect model performance, without hand-editing the treebank.

These roles are not mutually exclusive and are meant as indicative rather than exhaustive. The broader aim is to provide a transparent, modifiable pipeline that others can adapt to their own data and research questions, instead of a one-off script tied rigidly to a single release.

6.1.2. Packaging and License of the Converter

The converter is implemented in Python and distributed as a standalone codebase, separate from the treebank data itself. Its internal organization was described in Chapter 4; here the focus is on how it can be obtained and reused. From a user’s perspective, running the converter consists of invoking a single command with input and output file paths and, optionally, flags that control whether to perform only morphosyntactic conversion or the full pipeline.

To maximize reusability, the converter is released under the **MIT License**. This permissive license allows others to use, modify, and redistribute the code, including in proprietary or commercial projects, provided that the copyright notice and license text

are preserved. The choice of a permissive software license is intentional: the converter is meant to function as shared infrastructure that other projects can build on, fork, or integrate into their own tooling without licensing barriers.

It is important to stress that this software license applies only to the *converter code*. The MPDT-UD treebank released through the Universal Dependencies project (Chapter 5) remains distributed under the Creative Commons Attribution–ShareAlike 4.0 (CC BY-SA 4.0) license. Code and data are therefore governed by different, but compatible, licensing regimes: a permissive license for the software and a share-alike license for the annotated text.

6.1.3. Repository and Reproducible UD Integration

The converter is hosted in a public GitHub repository.¹⁰ The repository mirrors the modular structure described in Chapter 4 and provides the necessary entry points and documentation to run the pipeline on the MPDT source files. Keeping the code under version control has two practical consequences.

First, it supports **reproducibility**. The UD-conformant MPDT-UD 1.0 release corresponds to a specific state of the converter. By checking out the same version of the repository and re-running the pipeline on the same MPDT input, a researcher can reproduce the CoNLL-U file that underlies the treebank described in Chapter 5. Subsequent changes to the rules, or to the underlying MPDT annotation, can be tracked explicitly as commits, making it clear when and how the UD representation has evolved.

Second, it facilitates **integration with the UD ecosystem**. The converter is designed to fit into the standard workflow used for UD treebanks: it writes data in valid CoNLL-U format, is run together with the official `validate.py` script, and produces logs that make it easier to investigate validator messages and correct systematic issues. This alignment means that the same pipeline can be used both to generate the files that are submitted to Universal Dependencies and to produce local experimental variants for research.

In summary, the packaging and hosting of the converter aim to make it a maintainable, inspectable tool that can be reused beyond the scope of this thesis, while preserving a clear separation between the open-source software and the licensed treebank data it operates on.

6.2. Use Cases

The combination of a UD-conformant Middle Polish treebank and a reusable MPDT→MPDT-UD converter creates a small ecosystem rather than a single static

¹⁰Repository URL: <https://github.com/kvmilos/MPDT-to-UD-converter>.

dataset. The same pipeline that produced MPDT-UD 1.0 can be re-applied to future MPDT releases, and its rule-based design makes it possible to generate controlled variants of the data. This section sketches two broad families of use cases: studies in historical syntax and diachrony, and parser training and evaluation on diachronic data.

6.2.1. Historical Syntax and Diachrony

For historical linguists, MPDT-UD 1.0 offers a syntactically annotated sample of Middle Polish in a format that is comparable to modern Polish UD treebanks and to other languages. This makes a range of diachronic and typological questions more approachable in practice, because analyses can be carried out using standard UD tools and methods.

At a qualitative level, the treebank can be used to investigate phenomena described in Chapter 3 in a more systematic way: coordination patterns, the behaviour of certain conjunctions and other clause-linking expressions, the distribution of short-form adjectives and participles, and the persistence of dual number. The UD layer provides a normalized representation of these structures, which facilitates querying and comparison across texts, genres, and time periods.

At a quantitative level, the resource enables corpus-based studies that compare Middle Polish to modern Polish or to other Slavic languages within UD. Examples include: relative frequencies of non-projective dependencies; changes in the balance between analytic and synthetic expressions; shifts in the treatment of arguments and adjuncts; and diachronic changes in clause-linking strategies. Because the converter is rule-based and transparent, researchers can inspect or adjust specific mapping choices when such comparisons hinge on annotation details.

Finally, the availability of a UD-aligned treebank lowers the barrier to integrating Middle Polish into cross-linguistic projects that already rely on UD, such as typological surveys of coordination, argument structure, or word order. In such settings, MPDT-UD can function as a historical counterpart to modern Polish treebanks, allowing diachronically informed variants of existing studies without the need for bespoke tooling.

6.2.2. Parser Training and Evaluation

From an NLP perspective, MPDT-UD 1.0 and the converter together provide a testbed for working with historical Polish as an out-of-domain setting for dependency parsing. Because the treebank is encoded in standard CoNLL-U and passes the official UD validator, it can be used directly with existing UD parsers, taggers, and evaluation scripts.

One straightforward use case is to train a parser on modern Polish UD treebanks

and evaluate it on MPDT-UD 1.0, treating Middle Polish as a diachronic domain shift. Such experiments can illuminate which constructions cause systematic failures (e.g. non-projective configurations, complex predicate structures, or multiword connectives) and whether particular architectural or training choices mitigate these difficulties. Conversely, parsers can be trained on MPDT-UD itself to explore how far purely data-driven models can adapt to historical morphology and syntax given limited training material.

The converter also makes it possible to study the impact of annotation decisions on parsing performance. Because individual rule sets can be toggled or modified, researchers can generate alternative UD versions that differ in targeted ways—for example, different treatments of coordination heads, or alternative analyses of clause-linking constructions. Training and evaluating parsers on these variants would allow controlled investigations of how annotation schemes influence model accuracy and error profiles.

Finally, the pipeline can serve as a template for bringing other historical or specialized corpora into the UD ecosystem. Adapting the converter to a new source scheme (for instance, a related historical Polish resource or another Slavic treebank with PDB-inspired annotation) would enable comparable parser experiments across multiple time periods or varieties, using the same tooling and evaluation protocol.

6.3. Cognitive Science Perspective

Although this thesis is framed primarily as a contribution to corpus linguistics and language technology, its motivation and implications are closely connected to questions in cognitive science. Dependency treebanks are one of the main empirical interfaces between theories of mental representations of syntax, models of sentence processing, and observable patterns in text. By providing a UD-conformant treebank for Middle Polish, as well as a transparent conversion pipeline, this project creates a resource that can be used to investigate how processing constraints and category systems behave across time.

6.3.1. Processing Constraints

Psycholinguistic theories of sentence processing routinely appeal to structural constraints such as locality, memory limitations, and incremental interpretability. Many of these constraints can be operationalized directly on dependency structures: examples include measures of dependency length, the frequency and type of non-projective configurations, or the complexity of coordination and clause-linking patterns.

Because MPDT-UD 1.0 represents Middle Polish sentences in the same formalism as modern Polish UD treebanks, it becomes possible to compare such measures

diachronically. For instance, one can ask whether Middle Polish shows similar tendencies toward minimizing dependency length as modern Polish, or whether certain constructions—such as heavy preposed modifiers, deeply nested subordinate clauses, or elaborate coordination—were more frequent and thus may have imposed different processing demands on readers. The converter plays a central role here: it fixes a concrete mapping between MPDT and UD, allowing processing-oriented metrics to be computed in a consistent way across different stages of the language.

The treebank also facilitates the use of computational models as proxies for human processing. Modern parsers and language models, when trained or evaluated on MPDT-UD, can be treated as idealized processors exposed to Middle Polish input. Their error patterns and surprisals on particular constructions can then be related to hypotheses about human difficulty, ambiguity resolution, and expectation-based processing. In this way, a historically oriented treebank contributes to the broader cognitive-scientific programme of linking structural properties of language to processing behaviour.

6.3.2. Category Change Over Time

A second point of contact with cognitive science concerns how linguistic categories evolve over time. From a cognitive perspective, changes in part-of-speech inventories, morphological paradigms, or syntactic constructions are not purely formal: they reflect shifts in how speakers mentally organize their lexicon and grammar, and in how they generalize across usage patterns.

The conversion of MPDT to UD requires making explicit decisions about how Middle Polish forms relate to the category system used for contemporary languages in UD. Examples include mapping short-form adjectives and participles to appropriate UD categories, deciding how to analyse certain conjunctions and clause-linking expressions, or preserving dual number in the feature system even though it is lost in modern Polish. Each of these choices embodies a hypothesis about the continuity or reanalysis of categories across time: whether a given construction is best understood as an instance of a stable category (e.g. adjectives with changing surface forms) or as evidence of an ongoing shift.

Because the converter is rule-based and modifiable, it can support explicit experimentation with alternative categorization hypotheses. Researchers interested in diachronic category change can adjust the mapping rules—for instance, by treating certain borderline items as adpositions rather than adverbs, or by collapsing dual forms with plural—and then regenerate a full UD annotation under each hypothesis. Comparing the resulting distributions, as well as the behaviour of parsers or language models trained on them, offers one way of probing how different categorizations interact with usage patterns and processing models.

More broadly, situating Middle Polish within the UD inventory highlights the ten-

sion between language-specific categories and cross-linguistic, cognitively plausible abstractions. If historically attested forms in Middle Polish fit naturally into the same UD categories as their modern counterparts, this supports the view that at least some parts of the category system are cognitively robust over centuries. If they resist such alignment, this points to domains where grammaticalization, lexicalization, or shifts in constructional meaning have substantially reorganized the underlying mental representations. In this way, the technical work of designing and implementing the converter feeds back into core questions about how linguistic categories are represented and how they change over time in the minds of speakers.

6.4. Future Work

The converter and the initial MPDT-UD 1.0 release are intended as a starting point rather than a finished product. As MPDT grows and other historical resources become available, both the pipeline and the resulting treebank can be extended in several directions. This section highlights two main axes: expanding coverage and phenomena, and generalizing or automating parts of the conversion workflow.

6.4.1. Coverage and Phenomena

At present, MPDT-UD 1.0 mirrors the scope of MPDT as described in Chapter 3: prose, sentence length restricted to 10–50 tokens, and exclusion of poetry and sentences with extensive Latin insertions. As MPDT annotation progresses, a natural first direction for future work is simply *more data*: adding further texts, relaxing sentence-length constraints, and gradually incorporating more syntactically challenging material (possibly including poetry and mixed-language passages).

Each such extension would test the robustness of the converter and likely surface new constructions. In the current release, some validation issues were resolved as genuinely idiosyncratic cases—single sentences or very small clusters that did not justify complex general rules. With a substantially larger treebank, these *isolated edge cases* may turn out not to be isolated at all, but instances of broader patterns. In that scenario, the relevant ad-hoc fixes can be promoted to explicit, documented rules, or refactored into more general modules that handle an entire class of constructions. The existing logging infrastructure is well suited to this: it can be used to search for recurring change patterns and to identify where similar structures are still treated inconsistently.

A second line of work concerns broadening the range of phenomena encoded in the UD layer. For example, enhanced dependencies are currently focused on coordination and basic propagation of shared dependents; future versions could systematically capture additional phenomena.

6.4.2. Generalization and Automation

Beyond extending MPDT itself, the converter can be used as a template for bringing other corpora into UD. One mid-term goal is to factor out clearly which parts of the pipeline are specific to Middle Polish and which are reusable for other similar treebank schemes or related Slavic languages. This would make it easier to adapt the codebase to, for example, another historical Polish treebank or a contemporary resource that shares the same dependency labels but differs in detail.

A related direction involves adding more automation around rule design and maintenance. At the moment, conversion rules are handwritten and motivated by linguistic analysis and validator feedback. Future work could introduce lightweight tooling that helps discover candidate rules—for instance, by mining frequent configurations in the logs, comparing parser outputs on MPDT-UD with the original MPDT annotation, or highlighting clusters of residual validator warnings. Such tools would not replace expert judgment, but could shorten the loop between observing a systematic pattern and encoding it as an explicit rule.

Finally, the converter could be integrated more tightly into reproducible workflows used in NLP and cognitive science. Examples include packaging the pipeline as a Python library or command-line tool with versioned releases; providing ready-made configuration files for common scenarios (e.g. “tags-only conversion”, “full conversion with enhanced dependencies”); or building simple interfaces for inspecting sentence-level logs alongside converted trees. These steps would lower the barrier for other researchers to experiment with alternative mappings, extend the treebank, or port the approach to new data, while keeping the underlying principle of transparent, documented conversion intact.

Bibliography

- Brouwer, M., Brugman, H., and Kemps-Snijders, M. (2017). *MTAS: A Solr/Lucene based multi-tier annotation search solution*. CLARIN. URL: <http://www.ep.liu.se/ecp/136/002/ecp17136002.pdf>.
- Buchholz, S. and Marsi, E. (2006). “CoNLL-X Shared Task on Multilingual Dependency Parsing”. In: *Proceedings of the Tenth Conference on Computational Natural Language Learning (CoNLL-X)*. Ed. by Màrquez, L. and Klein, D. New York City: Association for Computational Linguistics, pp. 149–164. URL: <https://aclanthology.org/W06-2920/>.
- De Marneffe, M.-C., Manning, C. D., Nivre, J., and Zeman, D. (2021). “Universal Dependencies”. In: *Computational Linguistics* 47.2, pp. 255–308. ISSN: 0891-2017. DOI: 10.1162/coli_a_00402. URL: https://doi.org/10.1162/coli_a_00402.
- Gruszczyński, W., Adamiec, D., Bronikowska, R., Kieraś, W., Modrzejewski, E., Wieczorek, A., and Woliński, M. (2022). “The Electronic Corpus of 17th- and 18th-century Polish Texts”. In: *Language Resources and Evaluation* 56.1, pp. 309–332. ISSN: 1574-0218. DOI: 10.1007/s10579-021-09549-1. URL: <https://doi.org/10.1007/s10579-021-09549-1>.
- Nivre, J., de Marneffe, M.-C., Ginter, F., Hajič, J., Manning, C. D., Pyysalo, S., Schuster, S., Tyers, F., and Zeman, D. (2020). “Universal Dependencies v2: An Evergrowing Multilingual Treebank Collection”. In: *Proceedings of the Twelfth Language Resources and Evaluation Conference*. Ed. by Calzolari, N., Béchet, F., Blache, P., Choukri, K., Cieri, C., Declerck, T., Goggi, S., Isahara, H., Maegaard, B., Mariani, J., Mazo, H., Moreno, A., Odijk, J., and Piperidis, S. Marseille, France: European Language Resources Association, pp. 4034–4043. ISBN: 979-10-95546-34-4. URL: <https://aclanthology.org/2020.lrec-1.497/>.
- Przeiórkowski, A., Bańko, M., Górska, R. L., and Lewandowska-Tomaszczyk, B. (2012). *Narodowy Korpus Języka Polskiego*. Warszawa: Wydawnictwo Naukowe PWN. ISBN: 978-83-01-16700-4. URL: https://nkjp.pl/settings/papers/NKJP_ksiazka.pdf.
- Wieczorek, A. (2020). *Instrukcja anotowania drzew zależnościowych: Dodatek dla tekstów XVII- i XVIII-wiecznych*. Institute of Polish Language, Polish Academy of Sciences.

- Wieczorek, A. (2025). "Towards the Middle Polish Dependency Treebank". In: *Native Language in the 21st Century: System, Communication Practices and Education*. V & R Unipress.
- Wróblewska, A. (2018). "Extended and Enhanced Polish Dependency Bank in Universal Dependencies Format". In: *Proceedings of the Second Workshop on Universal Dependencies (UDW 2018)*. Ed. by de Marneffe, M.-C., Lynn, T., and Schuster, S. Brussels, Belgium: Association for Computational Linguistics, pp. 173–182. DOI: 10.18653/v1/W18-6020. URL: <https://aclanthology.org/W18-6020/>.
- Wróblewska, A. (2020). "Towards the Conversion of National Corpus of Polish to Universal Dependencies". In: *Proceedings of the Twelfth Language Resources and Evaluation Conference*. Ed. by Calzolari, N., Béchet, F., Blache, P., Choukri, K., Cieri, C., Declerck, T., Goggi, S., Isahara, H., Maegaard, B., Mariani, J., Mazo, H., Moreno, A., Odijk, J., and Piperidis, S. Marseille, France: European Language Resources Association, pp. 5308–5315. ISBN: 979-10-95546-34-4. URL: <https://aclanthology.org/2020.lrec-1.653/>.
- Wróblewska, A. (2023). *Instrukcja anotowania drzew w Polskim Banku Drzew Zależnościowych*.